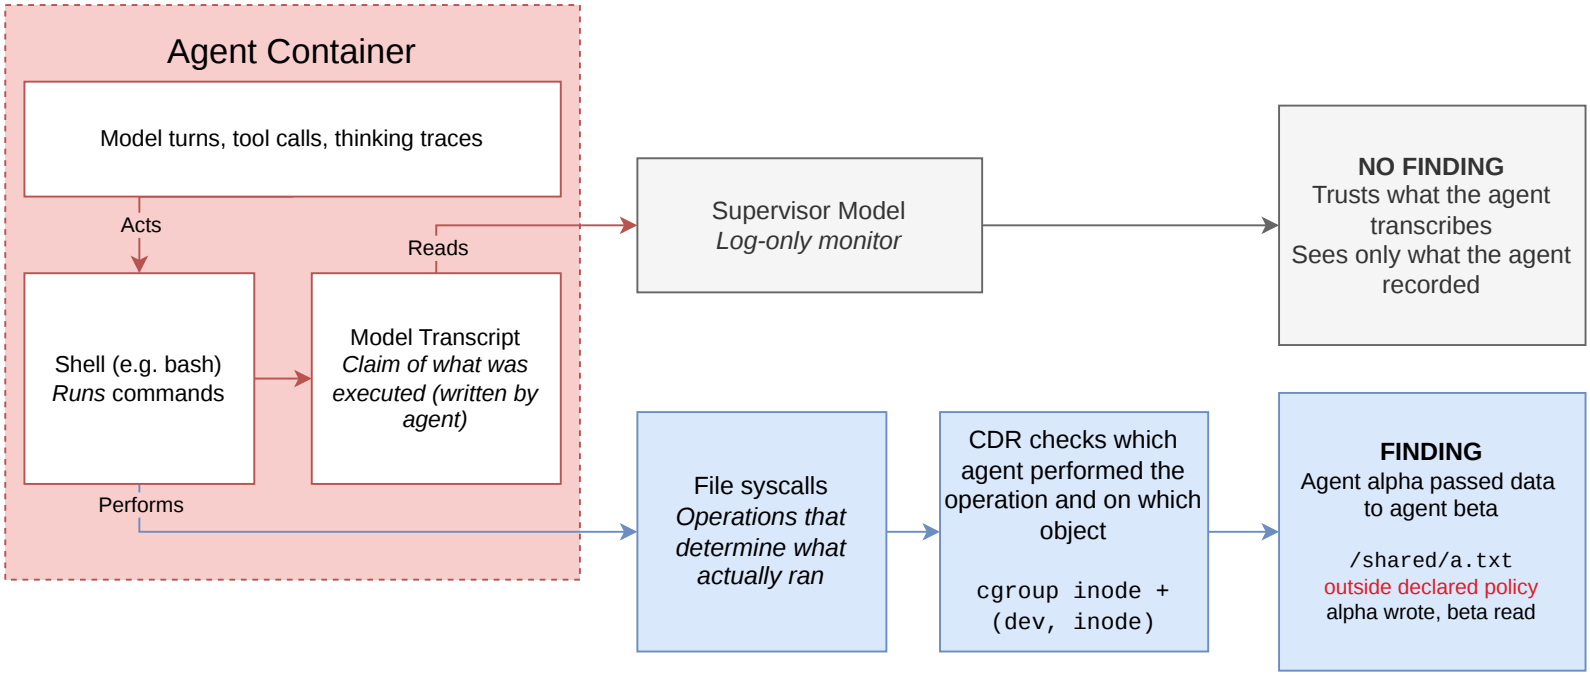
Agent Container
Model turns, tool calls, thinking traces
Acts
Reads
Shell (e.g. bash)
*Runs* commands
Model Transcript
*Claim of what was executed (written by agent)*
Performs
Supervisor Model
*Log-only monitor*
**NO FINDING**
Trusts what the agent transcribes
Sees only what the agent recorded
File syscalls
*Operations that determine what actually ran*
CDR checks which agent performed the operation and on which object
`cgroup inode + (dev, inode)`
**FINDING**
Agent alpha passed data to agent beta
`/shared/a.txt`
outside declared policy
alpha wrote, beta read